



Evaluating a Learning Algorithm

Bias vs. Variance

Review

Building a Spam Classifier



Video: Prioritizing What to Work On
9 min



Reading: Prioritizing What to Work On
3 min



Video: Error Analysis
13 min



Reading: Error Analysis
3 min

Handling Skewed Data

Using Large Data Sets

Review



Prioritizing What to Work On

System Design Example:

Given a data set of emails, we could construct a vector for each email. Each entry in this vector represents a word. The vector normally contains 10,000 to 50,000 entries gathered by finding the most frequently used words in our data set. If a word is to be found in the email, we would assign its respective entry a 1, else if it is not found, that entry would be a 0. Once we have all our x vectors ready, we train our algorithm and finally, we could use it to classify if an email is a spam or not.

Building a spam classifier

Supervised learning. x = features of email. y = spam (1) or not spam (0).
Features x : Choose 100 words indicative of spam/not spam.

E.g. deal, buy, discount, android, now, ...

$$x = \begin{bmatrix} 0 & 1 & 1 & 0 & \vdots & 1 & \vdots \\ \text{android} & \text{buy} & \text{deal} & \text{discount} & \vdots & \text{now} & \vdots \end{bmatrix} \quad x \in \mathbb{R}^{100}$$

$$x_j = \begin{cases} 1 & \text{if word } j \text{ appears in email} \\ 0 & \text{otherwise} \end{cases}$$

From: cheapsales@buystufffromme.com
To: ang@cs.stanford.edu
Subject: Buy now!

Deal of the week! Buy now!

So how could you spend your time to improve the accuracy of this classifier?

- Collect lots of data (for example "honeypot" project but doesn't always work)
- Develop sophisticated features (for example: using email header data in spam emails)
- Develop algorithms to process your input in different ways (recognizing misspellings in spam).

It is difficult to tell which of the options will be most helpful.